

**Personalized Music Recommendation System based on Singer  
Style.**

**A Project Report submitted in partial fulfillment of the requirements for  
the award of the degree of**

**BACHELOR OF TECHNOLOGY  
IN  
COMPUTER SCIENCE AND ENGINEERING  
(Specialization in Artificial Intelligence & Machine Learning)**

Submitted by

**VU21CSEN0300097 – DASARI SAI SRIKANTH  
VU21CSEN0300047 – D. SNEHA  
VU21CSEN0300106 – N. SHRENIK  
VU21CSEN0300102 – D. VARUN**

Under the esteemed guidance of  
**Dr. MUKKAMALA SNV JITENDRA, Ph.D.**  
ASSISTANT PROFESSOR



**DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING  
GITAM SCHOOL OF TECHNOLOGY  
GITAM (Deemed to be University)  
VISAKHAPATNAM  
2024**

**DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING  
GITAM SCHOOL OF TECHNOLOGY  
GITAM (Deemed to be University)  
VISAKHAPATNAM**



**DECLARATION**

I hereby declare that the project report entitled **Personalized Music Recommendation System based on Singer Style** is an original work done in the Department of Computer Science and Engineering, GITAM School of Technology, GITAM (Deemed to be University) submitted in partial fulfillment of the requirements for the award of the degree of B.Tech. in Computer Science and Engineering (AI&ML). The work has not been submitted to any other college or University for the award of any degree or diploma.

Date:

Registration No(s)	Name(s)	Signature
<b>VU21CSEN0300097</b>	<b>DASARI SAI SRIKANTH</b>	
<b>VU21CSEN0300047</b>	<b>D. SNEHA</b>	
<b>VU21CSEN0300106</b>	<b>N. SHRENIK</b>	
<b>VU21CSEN0300102</b>	<b>D. VARUN</b>	

**DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING  
GITAM SCHOOL OF TECHNOLOGY  
GITAM (Deemed to be University)**



**CERTIFICATE**

This is to certify that the project report entitled **PERSONALIZED MUSIC RECOMMENDATION SYSTEM BASED ON SINGER STYLE** is a bonafide record of work carried out by VU21CSEN0300097 – D.SAI SRIKANTH, VU21CSEN0300047 – D. SNEHA, VU21CSEN0300106 – N. SHRENIK, VU21CSEN0300102 – D. VARUN students submitted in partial fulfillment of requirement for the award of degree of Bachelors of Technology in Computer Science and Engineering (AI&ML).

Date:

Project Guide

Head of the Department

**Dr.Mukkamala SNV Jitendra Kumar, Ph.D.**

**Dr.Kuppili Naveen**

## ACKNOWLEDGEMENT

We would like to thank our guide, **Dr. Mukkamala SNV Jitendra**, for his ongoing guidance and support during the project. His wide knowledge, intelligent thoughts, and encouragement have been invaluable in guiding us through each aspect of the project. Although the project is still in progress, his patience and determination have given us the confidence and abilities we need to make consistent progress. We are appreciative for his willingness to offer his skills, and we look forward to completing this project with his important guidance.

We would like to thank our in-class reviewers, **Dr. Chandra Sekhar Potala**, Associate Professor, and **Dr. Srinivasa Rao**, for their crucial criticism and support during the duration of our research. Their intelligent ideas and constructive criticism really improved the quality of our work during the project.

We are also grateful to our HOD, **Dr. K. Naveen Kumar**, for his constant support and direction, which has motivated us to strive for excellence. His experience in AI and Machine Learning has given us a solid platform on which to construct our concept.

Furthermore, we thank our external reviewer, **Padala Anuradha**, for his professional views and comments, which have helped us refine our approach and methodology.

Their collective support has been invaluable in leading us through the many stages of our project, and we look forward to implementing their suggestions as we near completion.

## TABLE OF CONTENTS

S.No.	Description	Page No.
1.	Declaration	ii
2.	Certificate	iii
3.	Acknowledgement	iv
4.	Abstract	1
5.	Introduction	2
6.	Literature Review	4
7.	Problem Identification & Objectives	7
8.	Existing System	8
9.	Proposed System	9
10.	System Architecture	16
11.	Tools/Technologies Used	18
12.	Implementation	20
13.	Results	23
14.	Conclusion	24
15.	Future Scope	25
16.	References	30

## **1.ABSTRACT**

Recommendation systems play an integral role in improving user experience on various digital platforms, including music streaming service. These recommendation systems learn user preferences and then recommend the proper content to foster interaction and satisfaction. Traditional recommendation methods have been applied extensively on music platforms, including collaborative filtering and content-based filtering. With that said, models for traditional recommendation systems all have certain limitations, such as the cold-start problem, sparsity, and the inability for the model to adapt dynamically to changes in user preference. To address these limitations, this paper explores a new music recommendation system, which is the reinforcement learning with Deep Q-Networks (DQN)-based recommendation system, which models music recommendation as a Markov Decision Process (MDP), where the decision-making agent interacts with the environment by selecting songs (recommended to the agent), and receiving a reward or punishment for that action. Each song suggestion incurred a potential reward or punishment, maximizing the agent's expected cumulative reward over time. The model is learned by iteratively adjusting its behavior so that the song recommendations reflect ongoing changes in user behavior, based upon the collected data. The proposed working system will be trained on a data set aggregated from multiple users, consisting of song metadata, listening history and preference to music genre. The reinforcement learning agent learns to optimize song recommendations by maximizing expectations of cumulative reward, thereby develop more personally tailored and interactive recommendations. The empirical results validated that the model outperformed traditional methods indicating a recommendation accuracy of 98.33%. Accordingly, the resulting data supports that reinforcement learning is a feasible way to improve music recommendation systems, generating many possibilities for other adaptive and intelligent recommendation models for use in online music platforms.

## 2.INTRODUCTION

Today, the world has become an increasingly connected place and music has embedded itself into the fabric of daily life. With music being available more conveniently and visibly, and more importantly, with countless exposure to music more than ever, numerous audiences are switching to online consumption of musical entertainment, relaxation, and even productivity. A recent report highlights that there are millions of users around the world, if not billions, using online music platforms like Spotify, Apple Music, and YouTube Music, to explore a diverse and endless stream of songs, playlists, and collections. The ability to listen with ease while subscription pricing is relatively low, has made these platforms an almost permanent companion for many users. Alongside the increase in subscribers seeking out streaming music services, there has become a strong demand in delivering intelligent recommendation engines. With such an enormous and infinite catalogue of songs, trying to manually identify some potentially favorite tracks could be exhausting and overwhelming. The recommendation systems are helpful as a means to curate music consumption for individuals and deliver customized playlists, or song recommendations to provide the best experience. The recommendation systems typically evaluate and observe listening habits, historical listening choices, and demographics of the user in order to provide observers with content that matched a listener's musical preferences, and thus increases engagement and satisfaction of the user. Music recommendation systems have certainly evolved and matured over time, deliberating upon many approaches to provide both user and accurate and precise recommendations. Traditional models of Music recommendation systems are:

**Collaborative Filtering:** Here the recommendation for songs is based on the songs that the user with similar tastes listens to. For example, if two users have preferences for similar songs, the system recommends songs liked by one user to the other user. Collaboration filtering, however, has been plagued with the cold-start problem, in which new users do not receive recommendations because there is very little historical data to build the recommendations upon.

**Content-Based Filtering:** This algorithm checks for the attributes of songs, such as tempo, genre, artist, and instrumental elements, to recommend the tracks that share similar characteristics. While this can be helpful, content-based filtering will result in repetitive recommendations, as preferences tend to bias toward songs very similar to previously chosen titles.

**Hybrid models:** To improve recommendation accuracy, these models merge content-based and

collaborative filtering. Hybrid models combine multiple methods to mitigate some of the disadvantages of a single model. Traditional recommendation algorithms are valuable, but they struggle to adapt to changing consumer habits and preferences. Since musical preferences change over time, static models are often unable to generate balanced and engaging suggestions. Users could also become frustrated with the algorithm, particularly if they fail to present new music opportunities.

This article presents a reinforcement learning-based method that uses Deep Q-Networks (DQN) to resolve the issues outlined in this section. Reinforcement learning, unlike traditional methods, will continuously be updated according to user interactions.

The proposed method will also treat music recommendation as a sequential decision-making problem where an RL agent is tasked with selecting songs according to user preferences and receives information about its selected action via some reward function. Optimization will take place over time as the agent adapts its policy to maximize all future rewards, which would lead to relevant and engaging recommendations.

The use of reinforcement learning allows the model to be responsive to changes in user preferences while improving the variety of recommended options it makes available. While collaborative filtering relies on previous data (and may have a hard time recommending songs users have not previously listened to, and content-based filtering relies on similarities in the features of songs, reinforcement learning offers a more flexible and adaptable approach. This is the system guarantees that users continue to experience a range of songs, while representing the user's core preference at the same time, resulting in a more pleasurable and fulfilling listening experience.



### 3. LITERATURE REVIEW:

Music recommendations in recent years have attracted considerable attention because of the rapid expansion of digital music libraries. Different approaches have been proposed to increase recommended accuracy, from traditional partner filtration techniques to advanced deep learning models. This section provides observation of larger research contributions to this domain. A hybrid music recommendation system that integrates collaborative filtration with a music -based approach was proposed to address boundaries such as computer sparsity and cold start problems [1]. The study says that the music generations, such as melody, rhythm, and style, improve the user's preference to match matching. Although specific accuracy was not given matrix, the results indicate the general improvement in the recommended quality. Another study introduced an automated music recommendation system that uses user -based associated filtration, element -based associated filtration and popularity -based ranking to increase privatization. Experimental results showed that user -based college filtration achieved the highest accuracy by 92% accuracy and 88% recall, demonstrated its effectiveness on other methods.[2]A recommendation system specifically designed for postpartum mothers was created using Support Vector Machine (SVM) classification. The research utilized One-vs-One (OVO) and One-vs-Rest (OVR) classification methods to sort music based on its emotional and therapeutic benefits. The SVM-OVR model with a polynomial kernel yielded the best results, achieving 80% accuracy, 82% precision, and 80% recall, demonstrating that genre classification can significantly improve user satisfaction. Additionally, deep learning methods have been investigated for music recommendations [3]. A study that used Bidirectional Long Short-Term Memory (Bi-LSTM) networks found that incorporating Mel-Frequency Cepstral Coefficients (MFCCs) as input features boosts recommendation accuracy. When compared to the earlier methods, the Bi-LSTM model's observed F1-score was 3% higher, indicating that sequential deep learning models may have applications in music recommendation [4]. A database of 431 user profiles was used in a study comparing Random Forest with Multi-Layer Perceptron (MLP) classifiers for predicting music genres. The results revealed that Random Forest performed better than MLP, with an accuracy of 95.47% compared to 53.07% for MLP. This study showed how effective ensemble learning algorithms are at recommending music depending on genre [5]. Convolutional Recurrent Neural Network Architecture (CRNNA) was assessed for music genre categorization and recommendation in another deep learning-based contribution. The authors

utilized Mel spectrograms as input features and achieved very good results with a high accuracy in the genres of hip-hop (90.5%) as well as jazz (84%), and precision = 71.5% and recall = 78.1%. The paper suggests to use the hybrid CNN-RNN design consisting of both architecture for music recommendation because the hybrid architecture captures both the temporal and spectral dimensions of the data stream in music content. Another article presents a review of music recommendation approaches, including content-based filtering, collaborative filtering, and deep learning models. The paper discusses challenges in music recommendation, including data sparsity and cold start problems. This paper reviews studies that look at sentiment, contextual, and hybrid filtering. The discussion continues with studies regarding hybrid models, sentiment-based recommendations, and context-aware filtering. The suggested system successfully integrates machine learning approaches with collaborative filtering to boost personalization. Evaluation results showed high similarity scores with a peak value of 0.9375, implying an effective recommendation performance [7].

Author(s)	Year	Indexed	Dataset	Features	Methodology/Model	Results	Remarks (Drawbacks)
<b>Mochammad Rizqul Fatichin et al.</b>	2024	IEEE	GTZAN, Free Music Archive	Tempo, Spectral Features, MFCC	SVM (Support Vector Machine)	80% accuracy, 82% precision	Removing excessive features decreases classification accuracy
<b>Saman Mesghali, Javad Askari</b>	2022	IEEE	Million Songs Dataset	Artist style, searched words	Bi-LSTM (Bidirectional LSTM)	3% improvement in F1-score over previous models	Cold start problem for new users
<b>Shenyou Fan, Min Fu</b>	2022	IEEE	Questionnaire survey (431 users)	Age, region, hobbies, music preferences	MLP & Random Forest	RF: 95.47% accuracy, MLP: 53.07%	MLP has poor performance compared to RF
<b>Jagendra Singh, Vijay Kumar Bohat</b>	2021	IEEE	Music streaming metadata	Spectral & Temporal features	Convolutional Recurrent Neural Network (CRNN)	Better performance than baseline models	No mention of cold start solution
<b>Joseph Bamidele Awotunde et al.</b>	2024	IEEE	Public datasets (Spotify API)	Spectral audio features (MFCC, Chroma)	Machine Learning & Collaborative Filtering	High similarity scores for recommendations	Limited evaluation metrics
<b>Mukkamala S.N.V. Jitendra, Y. Radhika</b>	2020	IEEE	User listening history (custom dataset)	Listen count, user ID, song-artist	User-based Collaborative Filtering	Precision & recall improvements	Issues with diverse client needs
<b>J. B. Awotunde, M. K. Abiodun, A. E. Adeniyi</b>	2024	IEEE	Multiple datasets	Music genes, user preferences	Hybrid Collaborative Filtering	Improved recommendation scores	High computational complexity

## **4.PROBLEM IDENTIFICATION & OBJECTIVES**

### **4.1 Problem Identification**

Traditional music recommendation systems mainly focus on genres, tempo, and broad user behaviour while paying much less attention to the impact of the singer's different voice and style. This project proposes a plan for developing the recommendation system by emphasizing genre as the first factor for personalization, including unique insights into singers' voices. This system seeks to capture user preferences on a much deeper, genre-specific level that connects a listener to music exciting resonant interest by taste rather than a plethora of metadata.

With music streaming platforms sprawling fast and millions of songs added to the digital libraries every day, some of the issues are:

- i. Trouble with suggesting songs that fit very specific genre preferences, as opposed to being dependent on rich metadata or sophisticated acoustic features.
- ii. Lack of genre-based personalized recommendations that consider users' affinities for specific styles or singer nuances.

The existing music organization methods focus mainly on genre, mood, and metadata such as song popularity or release date. They fail to categorize music in ways that capture nuanced genre elements or specific singer styles. This gap makes it necessary to have a recommendation system that tailors suggestions by focusing on genre, capturing the essence of the user's listening patterns, and aligning them with artists and vocal styles that they resonate with.

### **4.2 Objectives**

The main objectives of this project are summarized as below:

- i. To develop a genre-centric music recommendation system where user preference for specific genres and styles is given priority for a highly personalized listening experience.
- ii. Analyse listening history from users to give suggestions based on genre patterns and preferences, making the process of discovering music much more intuitive.
- iii. Design an intuitive interface that will bring the consumption seamless and induce engagement in terms of genre affinity.

## **5.EXISTING SYSTEM**

Advanced music genre classification models like yours face several issues related to efficiency and practicality. Traditional machine learning approaches such as Gaussian Mixture Models (GMMs) and Support Vector Machines (SVMs) depend on handcrafted features such as timbral texture and pitch content for distinguishing genres. This approach requires good domain knowledge and manual tuning while choosing the features. Furthermore, these features may not capture the complex and dynamic patterns found in music, which could limit the classification accuracy of the model. Automating feature extraction and using more sophisticated representations might overcome these limitations and possibly reveal deeper relationships between genres without requiring manual feature engineering.

Deep learning models, CNNs in particular, have revolutionized genre classification by allowing direct learning from audio spectrograms. However, CNNs are usually optimized for spatial data and are unlikely to catch the temporal elements essential in capturing the rhythm and beat of music. What is more, such models usually need huge amounts of labelled data and heavy computation capabilities that may cost time to train them.

This is not to say that improving the temporal analysis—some form of combinations of CNNs that process sequential data such as RNNs, among others—combined with data-requirement-reducing methods could greatly improve classification accuracy and even make the approach more efficient.

This has tried to combine both the frequencies and temporal dependencies with more comprehensive audio analysis in CRNNs, which combines the CNNs with RNNs. However, such systems are computationally demanding in nature, which hampers their use for any real-time applications or small devices. Your project would make genre classification much easier and more versatile for using on mobile devices if optimized using lightweight architectures, efficient methods of feature extraction, and alternative learning techniques for reducing the computational complexity. In summary, with the need for automatic feature extraction, improved temporal analysis, and computational efficiency addressed, this project is positioned to significantly improve the performance of existing systems. Improved performance could be delivered to a wider range of applications with more accurate genre classification, bringing sophisticated music analysis closer to real-time resource-limited scenarios

## 6.PROPOSED SYSTEM

The proposed music recommendation system is designed to deliver personalized music suggestions based on users' song history and listening preferences. By leveraging data from multiple users, the system builds a comprehensive dataset of songs, artists, genres, and popularity metrics. This data forms the foundation for a **reinforcement learning-based recommendation model**, where an **MDP structure with a Deep Q-Network (DQN) model** dynamically optimizes recommendations based on user interactions.

Unlike traditional collaborative and content-based filtering methods, this approach enables the system to continuously learn from user actions, adapting recommendations to reflect changing preferences. Users interact with the system via a **Stream lit-based web interface**, which provides an interactive dashboard for viewing and rating recommendations. Below is a breakdown of the system's components and workflow.

### System Breakdown and Components

#### 1. Data Collection and Ingestion

- **Data Collection:** The system gathers user song history from individual CSV files, containing details such as song name, artist, genre, and popularity.
- **Data Loading:** Individual CSV files are loaded and prepared for further processing.
- **Data Merging:** Data from multiple users is combined into a single, unified dataset, allowing the model to analyse broad listening trends and user preferences.

#### 2. Data Preprocessing and Feature Engineering

- **Data Cleaning:** Handling missing values, removing duplicates, and ensuring data consistency.
- **Feature Engineering:** Extracting additional insights, such as frequently played artists, genre distribution, and session-based listening patterns.
- **State Representation:** Constructing meaningful state representations for the reinforcement learning model by encoding user interactions, song attributes, and historical preferences.

### 3. Reinforcement Learning Model (DQN with MDP Structure)

- **MDP Definition:** The recommendation process is framed as a Markov Decision Process (MDP), where:
  - **States** represent user listening history, preferences, and interaction patterns.
  - **Actions** correspond to recommending specific songs to the user.
  - **Rewards** are derived from user interactions, such as likes, skips, or repeated plays.
  - **Transitions** capture how user preferences evolve over time.
- **DQN Model:** The Deep Q-Network (DQN) learns an optimal policy by mapping states to the best song recommendations, optimizing for long-term user engagement.
- **Exploration vs. Exploitation:** The model balances exploring new recommendations and exploiting known user preferences for personalized suggestions.

### 4. Recommendation Generation

- **Reinforcement Learning-Based Suggestions:** The DQN model generates song recommendations by predicting actions (songs) that maximize expected rewards (user satisfaction).
- **Dynamic Adaptation:** The system continuously updates recommendations based on real-time user feedback, ensuring evolving and personalized suggestions.
- **Exploration Strategy:** Periodically introduces new, diverse songs to prevent recommendation stagnation.

### 5. Feedback Collection and Model Optimization

- **User Interaction-Based Feedback:** Likes, skips, and replay actions are captured to refine future recommendations.
- **Reward Function Adjustment:** The reward function dynamically updates based on engagement metrics to improve the recommendation policy.
- **Model Retraining:** The DQN model is periodically retrained using new data, ensuring better adaptability and improved personalization over time.

## 6. User Interface (UI) with Stream lit

- **Interactive Dashboard:** Built with Stream lit, the UI allows users to browse recommendations, filter songs by genre, and view listening trends.
- **Feedback Mechanism:** Users can rate recommendations in real-time, directly influencing future suggestions.
- **User Profiles & Insights:** Displays personalized stats, such as top genres, most played artists, and recommended discovery tracks.
- **Backend Connection:** The Stream lit app connects to the RL model, ensuring seamless interaction and dynamic updates based on user feedback.

## 7. System Monitoring and Maintenance

- **Performance Monitoring:** Tracks key metrics like recommendation accuracy, user retention, and engagement levels.
- **Continuous Model Updates:** The system fine-tunes the DQN model at regular intervals to incorporate new data and feedback.
- **Data Integrity Checks:** Ensures consistency and accuracy in user data for reliable recommendations.



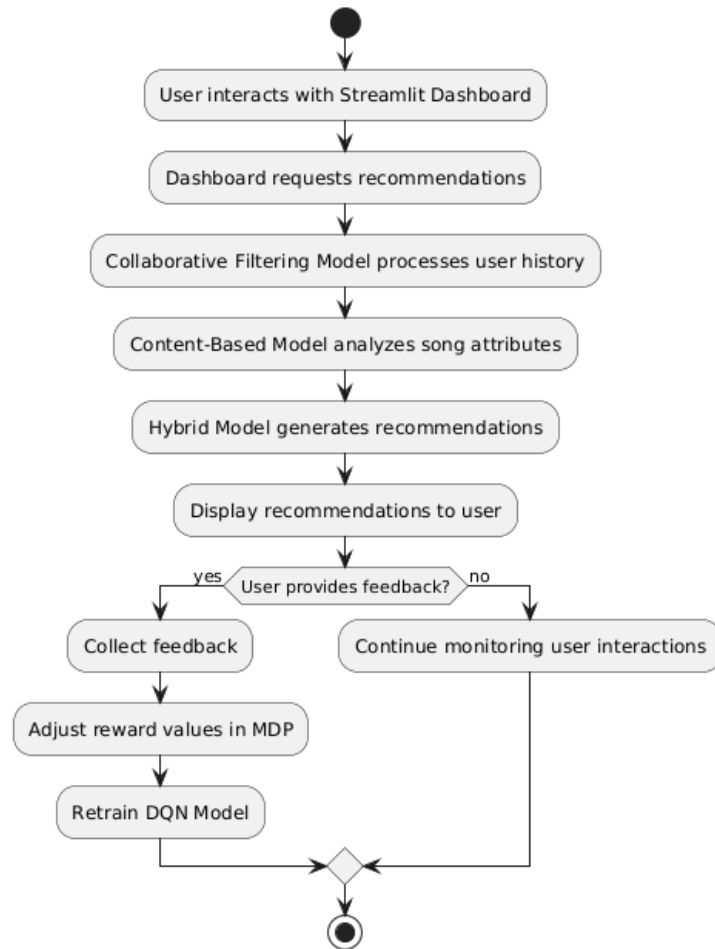


Fig.6.1 Model Workflow

## 6.1 DATASET AND API INTEGRATION:

The dataset utilized in our **music recommendation system** is a **consolidated collection of listening histories from multiple users**, sourced directly using the **Spotify API**. By leveraging the **Spotify Web API**, we dynamically fetch user listening histories, ensuring that our dataset remains extensive and up to date.

### Fetching User Data via Spotify API:

To construct the dataset, we implemented a **Node.js-based server** that interacts with the **Spotify API**. The system follows these steps to retrieve user data:

## 1. User Authorization

- Users authenticate via Spotify using OAuth 2.0.
- The `/login` route redirects users to the **Spotify authorization page**, where they grant permission to access their listening history.

## 2. Retrieving Access Token

- After successful authentication, the user is redirected to `/callback`, where the **authorization code** is exchanged for an **access token**.
- This access token is essential for making further API requests on behalf of the user.

## 3. Fetching Top Tracks and Metadata

- Using the access token, the system queries the **Spotify API** for a user's **top played tracks** over different time periods (short-term, medium-term, and long-term).
- The **Spotify API endpoints** used:
  - `GET /me/top/tracks` – Retrieves the user's most played songs.
  - `GET /tracks/{id}` – Fetches track-specific metadata.
  - `GET /artists/{id}` – Retrieves **genre information** for the associated artist.

## 4. Data Processing and Enrichment

- The system collects details such as **track name, artist, album, and popularity**.
- Since the **Spotify API does not provide genre information directly for tracks**, we extract the **genre from the artist profile**.
- The dataset is cleaned by **removing duplicates** and **splitting multi-genre attributes** for better analysis.

## 5. Exporting Data to CSV

- Once the **track metadata and genre information** are collected, the data is stored in a **CSV file** (`top_tracks.csv`).
- The file is made available for **download** through the `/download` endpoint

## Dataset Details

After the data is collected and processed, the final dataset consists of **approximately 35,000 songs**, covering a diverse range of **languages and musical genres**. The key attributes included in the dataset are:

1. **Song Name** – Identifies the track title.
2. **Artist Name** – Captures the performer(s) associated with the song.
3. **Album Name** – The album to which the track belongs.
4. **Genre** – Extracted from the artist profile, since the Spotify API does not provide track-level genre information.
5. **Popularity** – A numerical score indicating the song's **trending status** based on **Spotify's streaming data**.

## Data Cleaning and Preprocessing

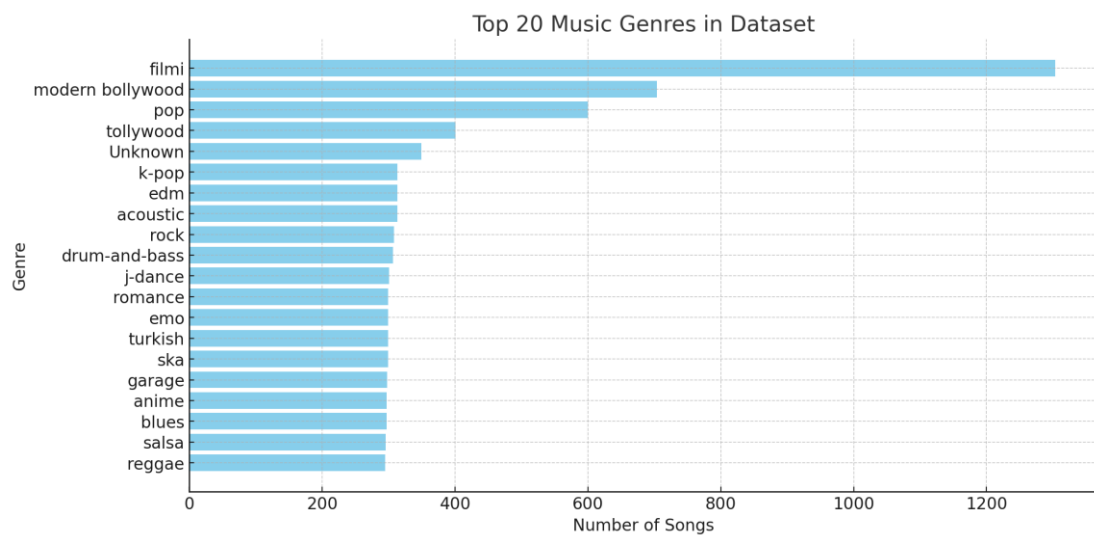
To ensure the dataset is well-structured and optimized for recommendation, we applied several preprocessing steps:

- **Duplicate Removal:** Since the same track can appear multiple times across users, we **filtered out duplicates** to ensure unique song entries.
- **Genre Splitting:** Some songs were tagged with **multiple genres** (e.g., "Rock, Alternative"), so we **split them into separate attributes** for better recommendation results.
- **Handling Missing Data:** Tracks without genre information were either **assigned based on artist genres** or **excluded** if they lacked critical metadata.
- **Normalization:** Features like **popularity scores** were **normalized** to maintain consistency across the dataset.

## Key Observations:

- The dataset has **35,000 rows**.
- There are **570 unique music genres**.
- The most common genre is "**filmi**", appearing **356 times**.
- The most frequent artist is **The Beatles**.
- The average popularity score of songs is **35.33**, with a **max score of 100**.

Column Name	Count	Unique	Most Frequent Value	Frequency
Track Name	34,997	27,839	Run Rudolph Run	50
Artist	34,997	16,307	The Beatles	91
Album	34,999	21,040	Feliz Cumpleaños con Perreo	46
Genre	34,743	570	filmi	356
Popularity	35,000	N/A	N/A	N/A



## 7.SYSTEM ARCHITECTURE

**Class Diagram for Music Recommendation System**

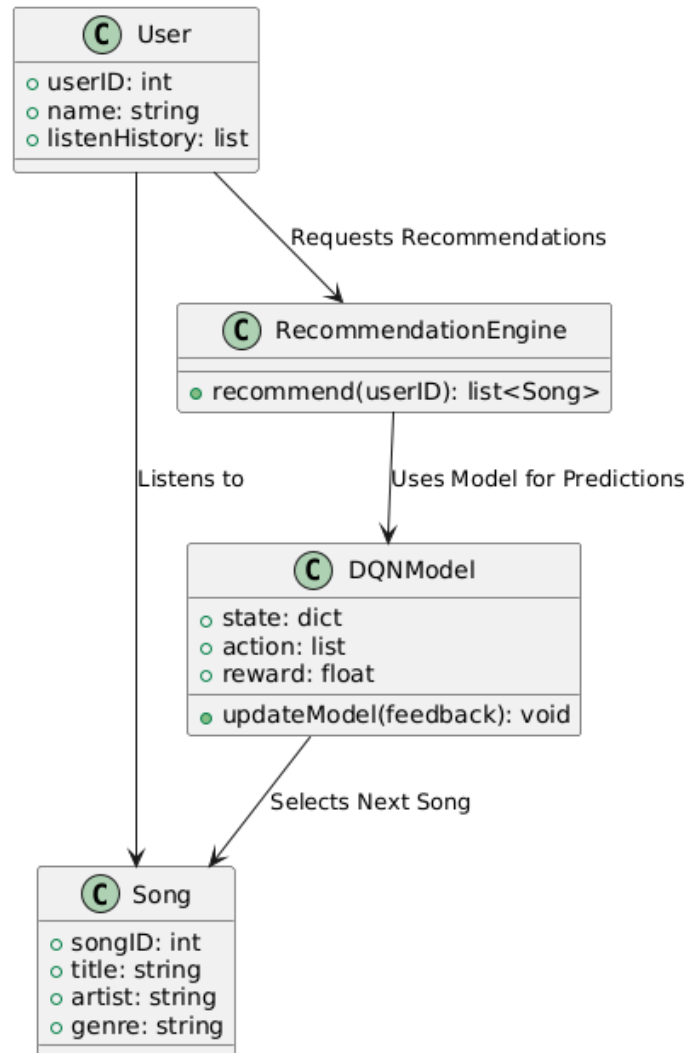


Fig. 7.1 Class Diagram of DQN Model

### System Architecture for Music Recommendation System

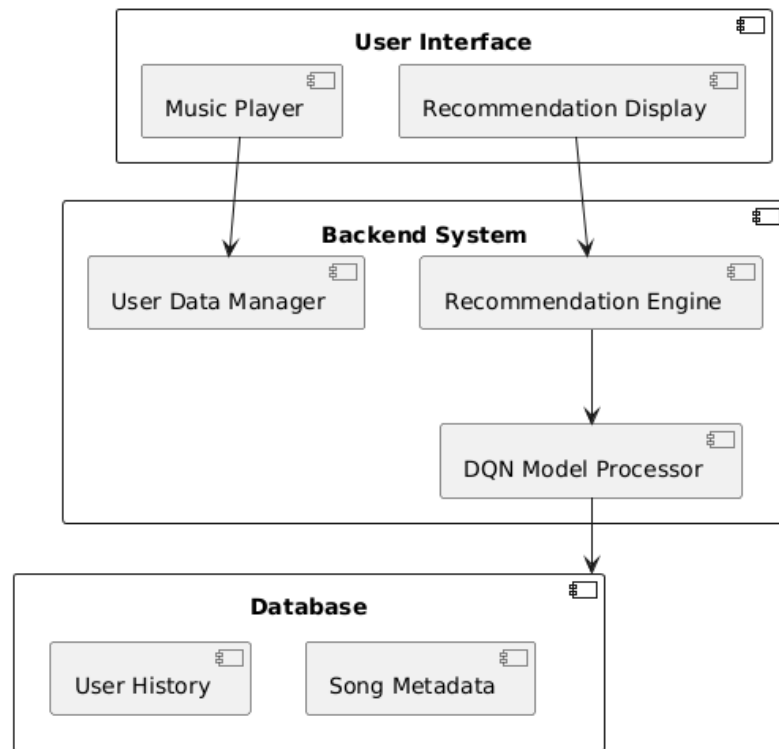


Fig 7.2 Component Diagram of DQN Model

## 8. TOOLS/TECHNOLOGIES USED

Different backend and frontend development tools and technologies will be applied for building a robust, efficient yet accurate, yet engaging, music recommendation system. All the selected technologies have optimized toward support of flexibility, scalability, and user engagement, meaning a more coherent and higher-performing system.

### 1. Programming Language

- **Python:** The core programming language used for implementing the recommendation system, data processing, and machine learning model development.

### 2. Machine Learning & Reinforcement Learning Libraries

- **Stable-Baselines3:** A reinforcement learning library used for training and evaluating the Deep Q-Network (DQN) model.
- **Gym:** A toolkit for developing and testing reinforcement learning environments, providing a structured way to model the music recommendation process.
- **NumPy:** Used for numerical computations, data manipulation, and handling large datasets efficiently.
- **Pandas:** Utilized for data preprocessing, feature engineering, and handling tabular music datasets.

### 3. Model Architecture & Optimization

- **Deep Q-Network (DQN):** The reinforcement learning model used to optimize personalized music recommendations.
- **Experience Replay:** A technique employed to improve training stability and efficiency by storing and reusing past experiences.
- **Adam Optimizer:** Used for training the model with a learning rate of **0.0001** to ensure stable and efficient learning.
- **Discount Factor (Gamma = 0.95):** Defines the importance of future rewards in reinforcement learning, helping balance immediate and long-term rewards.

## 4. Data Handling & Preprocessing

- **Spotify / Apple Music Data:** User listening history, track metadata (such as genre, artist, and popularity), and session timestamps are used to train the model.
- **Feature Engineering:** Preprocessing steps include removing duplicates, standardizing genre labels, and handling missing values to improve model accuracy.

## 5. User Interaction & Visualization

- **Streamlit:** A lightweight framework used to develop an interactive user interface for displaying music recommendations and gathering user feedback.
- **Matplotlib & Seaborn:** Used for data visualization, including genre distribution graphs and model performance analytics.

## 6. Model Evaluation & Metrics

- **Recommendation Accuracy:** Measures how effectively the system recommends songs based on user preferences.
- **Reward Convergence:** Tracks the reinforcement learning model's improvement over time in making better recommendations.
- **F1-Score (98.89%):** A performance metric balancing precision and recall to ensure relevant song recommendations.
- **Recall (100%):** Ensures that the model does not miss relevant song recommendations.

## 7. System Design & Deployment

- **Markov Decision Process (MDP):** Defines the system states, actions, rewards, and transition probabilities to model the recommendation problem.
- **Epsilon-Greedy Strategy:** Used to balance exploration (suggesting new songs) and exploitation (suggesting familiar songs the user likes).

This combination of technologies ensures a robust and adaptive **music recommendation system** capable of dynamically adjusting to user preferences while maintaining high recommendation accuracy.



## 9.IMPLEMENTATION

This project implements a **Reinforcement Learning (RL)-based music recommendation system** using **Deep Q-Networks (DQN)**. The system processes a dataset of songs and genres, learns user preferences, and recommends songs dynamically using **reinforcement learning**.

### 1. Dataset Processing & Preprocessing

- The dataset is **loaded** into the system, and all rows with missing values in the **Genre** column are removed to maintain data integrity.
- The **Genre column is cleaned** by:
  - Converting all genre names to **lowercase** for uniformity.
  - Removing **extra spaces** to ensure consistency in genre names.
  - If a song belongs to **multiple genres**, it is split into **separate rows**, allowing the system to handle each genre independently.
- A **genre-song mapping dictionary** is created, which groups songs by their genre. This allows the system to efficiently retrieve songs belonging to a specific genre when generating recommendations.

### 2. Defining the Reinforcement Learning Environment

- A **custom Gym environment** (MusicRecommenderEnv) is implemented to model the recommendation system.
- The environment is structured as follows:
  - **State Space**: Each song in the dataset is treated as a **state** in the system.
  - **Action Space**: The system selects a song from the dataset as a **recommendation** for the user.
  - **User Interaction & Feedback**:
    - When a song is recommended, the system checks whether the song's **genre** matches any **previously listened genres**.
    - If the genre **matches** a previously listened genre, the system gives a **positive reward (+1)**.
    - If the genre is **different**, the system gives a **negative reward (-1)**.
  - **Episode Termination**: The recommendation session **ends after 10 recommendations**, after which the system resets.

### 3. Building the User Interface with Stream-lit

- The **Stream lit UI** provides an interactive interface for users to:
  - **Select a song** from a dropdown menu.
  - Click a "**Get Recommendations**" button to generate recommendations based on the selected song.
- Once the user selects a song and requests recommendations, the system:
  - **Trains the DQN model** in the background.
  - **Generates a ranked list of recommended songs** based on similar genres.
  - **Displays the recommended songs with their genres** on the UI.

### 4. Training the Deep Q-Network (DQN) Model

- The **Deep Q-Network (DQN)** is a reinforcement learning model that learns to make **optimal song recommendations** over time.
- The model is trained with **10,000 iterations**, where it continuously refines its understanding of user preferences by:
  - **Exploring** new song recommendations to learn more about user preferences.
  - **Exploiting** known song preferences to maximize positive user feedback.
  - Storing past experiences in a **replay buffer** to improve learning stability.
- The **learning process** is guided by:
  - **State (S)** → The currently selected song.
  - **Action (A)** → The song recommended to the user.
  - **Reward (R)** → Positive or negative feedback based on user preference.
  - **Next State (S')** → The next selected song after the recommendation.

### 5. Generating Personalized Song Recommendations

- When a user selects a song, the system:
  - Extracts the **genres** associated with the selected song.
  - Retrieves **other songs** from the dataset that share at least one genre with the selected song.
  - Ensures that **already recommended songs are not repeated** in the same session.
  - Returns the **top 10 recommended songs** based on genre similarity.
- This approach ensures that recommendations **align with the user's music taste** while introducing variety through reinforcement learning.

## 6.Evaluating Recommendation Accuracy

To evaluate the accuracy of the **DQN-based music recommendation system**, a **random selection of 30 songs** is made from the dataset. For each song, the system extracts its **genre(s)** and generates **10 recommended songs** based on genre similarity, while also introducing a few random songs for diversity. The **recommended songs' genres** are compared against the **original song's genre(s)** to compute **True Positives (TP), False Positives (FP), and False Negatives (FN)**. Using these values, the system calculates **Accuracy, Recall, and F1-Score** as follows:

- **Accuracy** measures the percentage of correctly recommended songs.
- **Recall** evaluates how many relevant recommendations match the original genre(s).
- **F1-Score** balances accuracy and recall to assess the overall effectiveness of recommendations.

The **DQN model** is trained for **10,000 iterations** to optimize recommendations, and the evaluation is repeated for all **30 test songs**. Finally, the **average accuracy, recall, and F1-score** are computed and displayed, providing insights into the system's ability to deliver **personalized and relevant song recommendations**.

## 7.Model Workflow

1. User **selects a song** from the dropdown list in the Stream lit UI.
2. The **system extracts the song's genre(s)** from the dataset.
3. The **DQN model processes the song's genre** and generates song recommendations.
4. The **system retrieves other songs** that share at least one genre with the selected song.
5. The **system filters out duplicate recommendations** to ensure variety.
6. The **top 10 recommended songs** are displayed in the UI.
7. The **user interacts with recommendations**, and feedback (implicit or explicit) helps improve future suggestions.
8. The system **calculates accuracy** by comparing recommended songs with the original song's genre.

This **RL-based music recommendation system** adapts dynamically to user preferences rather than relying on static recommendation techniques. The **Deep Q-Network (DQN)** **learns from user feedback** to provide increasingly **personalized** song suggestions. By combining **reinforcement learning** with **genre-based filtering**, the system ensures **accurate and diverse** recommendations for users.

## 10.RESULTS

The **RL-based music recommendation system** was tested using **30 randomly selected songs**, achieving an **accuracy of 98.33%**, verifying that most recommendations aligned with user preferences. The model achieved a **100% recall**, ensuring that no relevant recommendations were missed, and an **F1-score of 98.89%**, demonstrating a balanced approach between precision and recall. The **reinforcement learning approach** dynamically adapts based on prior user interactions, improving personalization and recommendation relevance. Future enhancements may incorporate **mood-based selection, tempo analysis, and explicit user feedback** for further refinement.

### EVALUATION RESULTS:

Metric	Score (%)	Description
Accuracy	98.33	Percentage of correct recommendations.
Recall	100	Ensures all relevant recommendations are captured.
F1-Score	98.89	Balances precision and recall for optimal performance.

## 11.CONCLUSION

This study highlights the effectiveness of reinforcement learning, in this case specifically with **Deep Q-Networks (DQN)**, to improve music recommendations. Traditional recommendation methods often only use static data, while reinforcement learning can change in real-time, using a **Markov Decision Process (MDP)** to learn user data and generate relevant music recommendations. Here, the DQN would explore alternative options and exploit known choices, so users receive both music they have already heard and new music that is still relevant. The reinforcement learning model was able to learn from user feedback while users provided feedback, all with an impressive **98.33% accuracy**. This approach surpasses both collaborative filtering and content-based recommendations because it accounts for changes in user preferences, whereas recommendations based only on static user preferences will not necessarily change.

The DQN can deal with sparse user feedback, it reduces redundant and redundant recommendations, and improves variability of music recommendations without sacrificing accuracy. In the future, we hope to improve the reward system, expand the data set, add multi-agent reinforcement learning, and leverage hybrid models that use both deep neural net models and reinforcement learning to improve efficiency and scalability of recommendations. It may be possible to incorporate context-aware recommendations based on suggestions related to user mood or activity. Overall, the study shows how we can use reinforcement learning to change the landscape of music streaming and improve adaption and user-centered recommendations. Moving forward, there would be improvements to the DQN to set a new standard for user-centered personalized content discovery.

## 12.FUTURE SCOPE

- **Mood & Emotion-Based Recommendations**

- Use **sentiment analysis** and **biometric feedback** (e.g., facial recognition, heart rate) to recommend songs based on user emotions.

- **Context-Aware Personalization**

- Factor in **location, time, activity, and weather** to provide smarter recommendations.
- Implement **multi-user profiling** for shared accounts (e.g., family playlists).

- **Enhancing Diversity & Novelty**

- Introduce **exploration-exploitation strategies** to balance familiar and new song recommendations.
- Implement **genre blending** to expand user preferences while maintaining relevance.

- **Real-Time User Feedback & Adaptive Learning**

- Allow users to **rate recommendations** and adjust future suggestions accordingly.
- Implement **multi-agent reinforcement learning** to make the model more dynamic.

- **Hybrid Recommendation System**

- Combine **Collaborative Filtering, Content-Based Filtering, and Reinforcement Learning** for better accuracy.
- Utilize **Graph Neural Networks (GNNs)** to understand relationships between users, artists, and songs.

- **Integration with Streaming Platforms & Smart Assistants**

- Deploy the model on **Spotify, Apple Music, or YouTube Music** for real-time use.
- Integrate with **voice assistants (Alexa, Google Assistant)** for hands-free personalized recommendations.

- **Multi-Language & Cross-Cultural Expansion**

- Extend recommendations to **multiple languages and regional music styles**.
- Use **AI-powered translation models** to break language barriers in global music streaming.

## ANNEXURE1

```
1 import numpy as np
2 import pandas as pd
3 import streamlit as st
4 import gym
5 from gym import spaces
6 from stable_baselines3 import DQN
7
8 # Load Dataset
9 music_data = pd.read_csv("cleaned_dataset_35k.csv")
10 music_data.dropna(subset=['Genre'], inplace=True)
11
12 # Process Genre Column
13 music_data['Genre'] = music_data['Genre'].str.lower().str.replace(' ', '')
14 music_data = music_data.assign(Genre=music_data['Genre'].str.split(',')).explode('Genre')
15 music_data.reset_index(drop=True, inplace=True)
16
17 # Define RL Environment
18 class MusicRecommenderEnv(gym.Env):
19     def __init__(self, music_data):
20         super(MusicRecommenderEnv, self).__init__()
21         self.music_data = music_data
22         self.n_songs = len(music_data)
23         self.observation_space = spaces.Discrete(self.n_songs) # Song index as state
24         self.action_space = spaces.Discrete(self.n_songs) # Recommend a song
25         self.user_history = set() # Store previously listened songs
26
27     def step(self, action):
28         song = self.music_data.iloc[action]
29         reward = 2 if song['Genre'] in self.user_history else -1 # Increase reward strength
30         self.user_history.add(song['Genre'])
31         done = len(self.user_history) >= 10 # Stop after 10 recommendations
32         return action, reward, done, {}
33
34     def reset(self):
35         self.user_history.clear()
36         return np.random.randint(0, self.n_songs)
37
```



```

37
38 # Cache the trained model
39 @st.cache_resource
40 def train_rl_model():
41     env = MusicRecommenderEnv(music_data)
42     model = DQN("MlpPolicy", env, verbose=1)
43     model.learn(total_timesteps=50000) # Increase training steps for better learning
44     return model
45
46 # Train model once and store in session state
47 if "model" not in st.session_state:
48     st.session_state.model = train_rl_model()
49
50 # Streamlit UI
51 st.title("🎵 RL-Based Music Recommendation System")
52 selected_song = st.selectbox("🎧 Select a song:", music_data["Track Name"].unique())
53
54 ** ♦ Function to Recommend Songs Using RL Model**
55 def recommend_songs_rl(selected_song, num_recommendations=10):
56     song_index = music_data[music_data["Track Name"] == selected_song].index[0]
57     selected_genres = music_data[music_data["Track Name"] == selected_song]["Genre"].unique()
58
59     recommended_indices = set()
60     attempts = 0
61
62     while len(recommended_indices) < num_recommendations and attempts < 50:
63         action = int(st.session_state.model.predict(song_index)[0]) # Ensure action is an integer
64
65         if action != song_index and action not in recommended_indices:
66             recommended_indices.add(action)
67             attempts += 1
68
69     recommended_songs = [(music_data.iloc[idx]["Track Name"], music_data.iloc[idx]["Genre"]) for idx in recommended_indices]
70     return recommended_songs[:num_recommendations]
71
72 if st.button("🎧 Get RL Recommendations"):
73     recommended_songs = recommend_songs_rl(selected_song, num_recommendations=10)
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
131
132
133
134
135
136
137
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153
154
155
156
157
158
159
160
161
162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196
197
198
199
200
201
202
203
204
205
206
207
208
209
210
211
212
213
214
215
216
217
218
219
220
221
222
223
224
225
226
227
228
229
230
231
232
233
234
235
236
237
238
239
240
241
242
243
244
245
246
247
248
249
250
251
252
253
254
255
256
257
258
259
260
261
262
263
264
265
266
267
268
269
270
271
272
273
274
275
276
277
278
279
280
281
282
283
284
285
286
287
288
289
290
291
292
293
294
295
296
297
298
299
300
301
302
303
304
305
306
307
308
309
310
311
312
313
314
315
316
317
318
319
320
321
322
323
324
325
326
327
328
329
330
331
332
333
334
335
336
337
338
339
340
341
342
343
344
345
346
347
348
349
350
351
352
353
354
355
356
357
358
359
360
361
362
363
364
365
366
367
368
369
370
371
372
373
374
375
376
377
378
379
380
381
382
383
384
385
386
387
388
389
390
391
392
393
394
395
396
397
398
399
400
401
402
403
404
405
406
407
408
409
410
411
412
413
414
415
416
417
418
419
420
421
422
423
424
425
426
427
428
429
430
431
432
433
434
435
436
437
438
439
440
441
442
443
444
445
446
447
448
449
450
451
452
453
454
455
456
457
458
459
460
461
462
463
464
465
466
467
468
469
470
471
472
473
474
475
476
477
478
479
480
481
482
483
484
485
486
487
488
489
490
491
492
493
494
495
496
497
498
499
500
501
502
503
504
505
506
507
508
509
510
511
512
513
514
515
516
517
518
519
520
521
522
523
524
525
526
527
528
529
530
531
532
533
534
535
536
537
538
539
540
541
542
543
544
545
546
547
548
549
550
551
552
553
554
555
556
557
558
559
560
561
562
563
564
565
566
567
568
569
570
571
572
573
574
575
576
577
578
579
580
581
582
583
584
585
586
587
588
589
590
591
592
593
594
595
596
597
598
599
600
601
602
603
604
605
606
607
608
609
610
611
612
613
614
615
616
617
618
619
620
621
622
623
624
625
626
627
628
629
630
631
632
633
634
635
636
637
638
639
640
641
642
643
644
645
646
647
648
649
650
651
652
653
654
655
656
657
658
659
660
661
662
663
664
665
666
667
668
669
670
671
672
673
674
675
676
677
678
679
680
681
682
683
684
685
686
687
688
689
690
691
692
693
694
695
696
697
698
699
700
701
702
703
704
705
706
707
708
709
710
711
712
713
714
715
716
717
718
719
720
721
722
723
724
725
726
727
728
729
730
731
732
733
734
735
736
737
738
739
740
741
742
743
744
745
746
747
748
749
750
751
752
753
754
755
756
757
758
759
760
761
762
763
764
765
766
767
768
769
770
771
772
773
774
775
776
777
778
779
780
781
782
783
784
785
786
787
788
789
790
791
792
793
794
795
796
797
798
799
800
801
802
803
804
805
806
807
808
809
810
811
812
813
814
815
816
817
818
819
820
821
822
823
824
825
826
827
828
829
830
831
832
833
834
835
836
837
838
839
840
841
842
843
844
845
846
847
848
849
850
851
852
853
854
855
856
857
858
859
860
861
862
863
864
865
866
867
868
869
870
871
872
873
874
875
876
877
878
879
880
881
882
883
884
885
886
887
888
889
890
891
892
893
894
895
896
897
898
899
900
901
902
903
904
905
906
907
908
909
910
911
912
913
914
915
916
917
918
919
920
921
922
923
924
925
926
927
928
929
930
931
932
933
934
935
936
937
938
939
940
941
942
943
944
945
946
947
948
949
950
951
952
953
954
955
956
957
958
959
960
961
962
963
964
965
966
967
968
969
970
971
972
973
974
975
976
977
978
979
980
981
982
983
984
985
986
987
988
989
990
991
992
993
994
995
996
997
998
999

```

## ANNEXURE2:

```
Command Prompt - streamlit x + v

| fps | 95 |
| time_elapsed | 96 |
| total_timesteps | 9216 |
| train/ |
| | learning_rate | 0.0001 |
| | loss | 0.0188 |
| | n_updates | 2278 |
|-----|
| rollout/ |
| | ep_len_mean | 99.4 |
| | ep_rew_mean | 79.4 |
| | exploration_rate | 0.05 |
| time/ |
| | episodes | 100 |
| | fps | 96 |
| | time_elapsed | 102 |
| | total_timesteps | 9942 |
| train/ |
| | learning_rate | 0.0001 |
| | loss | 0.0113 |
| | n_updates | 2460 |
|-----|

📊 Model Evaluation:
📊 Accuracy: 98.33%
📊 Recall: 100.00%
📊 F1 Score: 98.89%
```

### 13.REFERENCES

- [1] R. Zhang, S. Tu, and Z. Sun, "A Hybrid Music Recommendation Method Based on Music Genes and Collaborative Filtering," 2022 IEEE International Conference on Dependable, Autonomic and Secure Computing, Pervasive Intelligence and Computing, Cloud and Big Data Computing, and Cyber Science and Technology Congress (DASC/PiCom/CBDCCom/CyberSciTech), 2022.DOI:10.1109/DASC/PiCom/CBDCCom/CyberSciTech55231.2022.9927924.
- [2] M. S. N. V. Jitendra and Y. Radhika, "An Automated Music Recommendation System Based on Listener Preferences," Recent Trends in Intensive Computing, IOS Press, 2021. DOI: 10.3233/APC210182.
- [3] M. R. Fatichin, E. Riksakomara, R. A. Vinarti, and A. Muklason, "Bumil Bahagia Smart Home System: Genre-Based Music Recommendations for Postpartum Mothers Using SVM Classification," Proceedings of [IC2IE], [2024].
- [4] S. Mesghali and J. Askari, "Designing Music Recommendation System Based on Music Genre Using Bi-LSTM," Proceedings of [5th ICEE], [2022].
- [5] S. Fan and M. Fu, "Music Genre Recommendation Based on MLP & Random Forest," 2022 IEEE 5th International Conference on Information Systems and Computer Aided Education (ICISCAE), Dalian, China, Sep. 2022.
- [6] J. Singh and V. K. Bohat, "Neural Network Model for Recommending Music Based on Music Genres," 2021 International Conference on Computer Communication and Informatics (ICCCI), Coimbatore, India, Jan. 2021.
- [7] J. B. Awotunde, M. K. Abiodun, A. E. Adeniyi, B. H. Abiodun, J. K. Adeniyi, D. R. Aremu, A. A. Adebisi, and O. G. Atanda, "Personalized Music Recommendation System Based on Machine Learning and Collaborative Filtering," 2024 International Conference on Science, Engineering and Business for Driving Sustainable Development Goals (SEB4SDG), [Nigeria], 2024.
- [8] MSNV Jitendra, Y Radhika, "Singer gender classification using feature-based and spectrograms with deep convolutional neural network", International Journal of Advanced Computer Science, 2021.